

AI-NATIVE B2B SALES OPERATIONS OS

A human-governed sales agent that connects customer channels, CRM memory, deterministic commercial logic, company knowledge and approval-controlled automation.

WORKING LOCAL PROTOTYPE

B2B MANUFACTURING


HUMAN-GOVERNED AI

Architecture and engineering brief

Designed and implemented as a local proof of architecture; production hardening and deployment are the next stage.


System architecture

The LLM interprets language; deterministic services and policies control business outcomes.



End-to-end sales workflow

One persistent thread follows the customer from first inquiry to operational handoff.



Design rule

Commercial interest is not treated as an order. A deal becomes won only after valid acceptance, PO validation, revalidation and controlled order creation.

Why this is more than an LLM chatbot

The architecture separates language intelligence from enforceable commercial authority.

Concern	Conventional AI assistant	This architecture
Prices and margins	May infer values from prompts or documents	Exact PostgreSQL master data and deterministic calculation
Business authority	Prompt instructions become the practical boundary	Separate AI principal, scopes, policies and approval gates
Customer memory	Conversation history is fragmented	Permanent customer ID, interactions, Customer 360 and semantic notes
Reliability	A repeated request can repeat the side effect	Idempotency keys, durable jobs, checkpoints and uniqueness constraints
Explainability	A generated answer is difficult to audit	Policy version, facts, evidence, actor, timestamps and execution result
Human control	Human review is informal	Explicit pending approvals that expire and can be consumed once

Core engineering principle

Gemini determines what a person means. PostgreSQL determines the exact business facts. The authority engine determines what Jarvis may do. A human approves exceptions.

4 Current maturity and opportunity

A tested architectural prototype with a clear path to production.

Implemented in the local prototype

- Website, email and WhatsApp ingestion architecture
- Persistent LangGraph inquiry-to-handoff workflow
- PostgreSQL CRM, customer identity and Customer 360
- Structured catalog, feasibility and pricing repositories
- Qdrant RAG for company knowledge and semantic memory
- Quotations, follow-ups, negotiation, PO and handoff
- Founder CRM, dashboard, tasks and approvals
- Controlled Jarvis tools and deterministic authority engine
- Tenant isolation, idempotency, auditing and regression tests

Next production milestones

- Batch 4: complete AI action-request lifecycle ledger
- Batch 5: conversational founder-facing Jarvis orchestrator
- Production deployment, observability and security hardening
- Server Qdrant, secrets management and enterprise identity
- ERP integrations such as SAP, Tally, Odoo or Zoho
- Load, failure-recovery and provider integration testing
- Predictive reorder and recommendation services

30-second technical pitch

I am building an AI-native B2B Sales Operations OS that moves an inquiry through customer identification, feasibility, pricing, quotation, follow-up, negotiation, PO validation and operational handoff. The LLM handles interpretation and communication, while exact data, deterministic policies, approvals, idempotency and auditing govern every business action. I am looking for an opportunity to develop this architecture with experienced AI builders and take it from a tested local prototype to a secure production platform.